

# Tutorial: Large-Scale Graph Processing in Shared Memory

---



Julian Shun (University of California, Berkeley)  
Laxman Dhulipala (Carnegie Mellon University)  
Guy Blelloch (Carnegie Mellon University)

# Tutorial Agenda

- 2:00-3:30pm
  - Overview of graph processing and Ligra
  - Walk through installation
  - Do examples in Ligra
- 3:30-4:00pm
  - Break
- 4:00-5:30pm
  - Implementation details of Ligra
  - Overview of other graph processing systems
  - Exercises

Slides available at <https://github.com/jshun/ligra/tree/master/tutorial/tutorial.pdf>

# Graphs are everywhere!



- Can contain billions of vertices and edges!



6.6 billion edges

Common Crawl


128 billion edges

facebook

~1 trillion edges [VLDB 2015]

# Graph processing challenges

- Many random memory accesses, poor locality
  - Relatively high communication-to-computation ratio
  - Varying parallelism throughout execution
  - Race conditions, load balancing
- 
- Need to efficiently analyze large graphs
    - Running time efficiency
    - Space efficiency
    - Programming efficiency



# Ligra Graph Processing Framework

EdgeMap

VertexMap

Breadth-first search  
Betweenness centrality  
Connected components  
K-core decomposition  
Belief propagation  
Maximal independent set  
...

Single-source shortest paths  
Eccentricity estimation  
(Personalized) PageRank  
Local graph clustering  
Biconnected components  
Collaborative filtering  
...

*Simplicity, Performance, Scalability*

# Graph Processing Systems

- Existing (at the time Ligra was developed): Pregel/Giraph/GPS, GraphLab, Pegasus, Knowledge Discovery Toolbox, GraphChi, and many others...
- Our system: Ligra - Lightweight graph processing system for shared memory

*Takes advantage of “frontier-based” nature of many algorithms (active set is dynamic and often small)*



# Steps for Graph Traversal

Many graph traversal algorithms do this!

Graph

- Operate on a subset of vertices
- Map computation over subset of edges **in parallel**
- Return new subset of vertices
- Map computation over subset of vertices **in parallel**


VertexSubset

EdgeMap

VertexMap

*We built the **Ligra** abstraction for these kinds of computations*

*Think with flat data-parallel operators*



Intel®  
**Cilk™ Plus**

**OpenMP**


Shared memory  
parallelism

Optimizations:

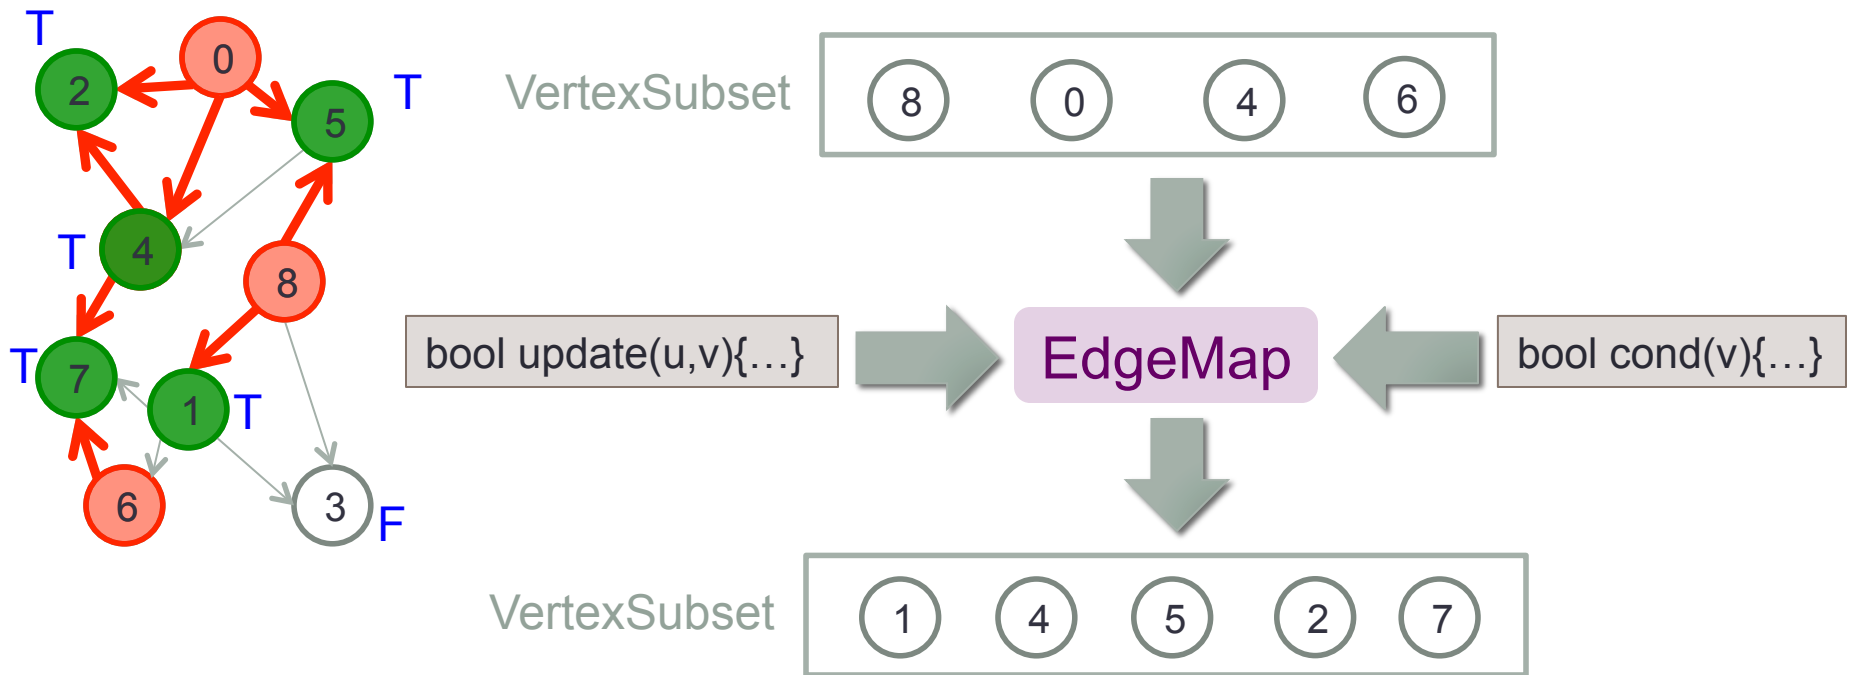
- hybrid traversal
- graph compression



# Ligra Framework



# Ligra Framework




# Breadth-first Search in Ligra

```
parents = {-1, ..., -1}; // -1 indicates "unexplored"
```

```
procedure UPDATE(s, d):  
    return compare_and_swap(parents[d], -1, s);
```

```
procedure COND(v):  
    return parents[v] == -1; // checks if "unexplored"
```

```
procedure BFS(G, r):  
    parents[r] = r;  
    frontier = {r}; // VertexSubset  
    while (size(frontier) > 0):  
        frontier = EDGEMAP(G, frontier, UPDATE, COND);
```




# Install and code examples in Ligra

# Breadth-first Search (BFS)


- Compute a BFS tree rooted at source  $r$  containing all vertices reachable from  $r$

Frontier




# Connected Components

- Takes an unweighted, undirected graph  $G = (V, E)$
- Returns a label array  $L$  such that  $L[v] = L[w]$  if and only if  $v$  is connected to  $w$



# Parallel Label Propagation Algorithm




- Processing all vertices in each iteration seems wasteful
  - Optimization: only place vertices who changed on frontier
- Warning: this algorithm is only good for low-diameter graphs

# Single-Source Shortest Paths

- Takes a weighted graph  $G = (V, E, w(E))$  and starting vertex  $r \in V$
- Returns a shortest paths array  $SP$  where  $SP[v]$  stores the shortest path distance from  $r$  to  $v$  ( $\infty$  if  $v$  unreachable from  $r$ )

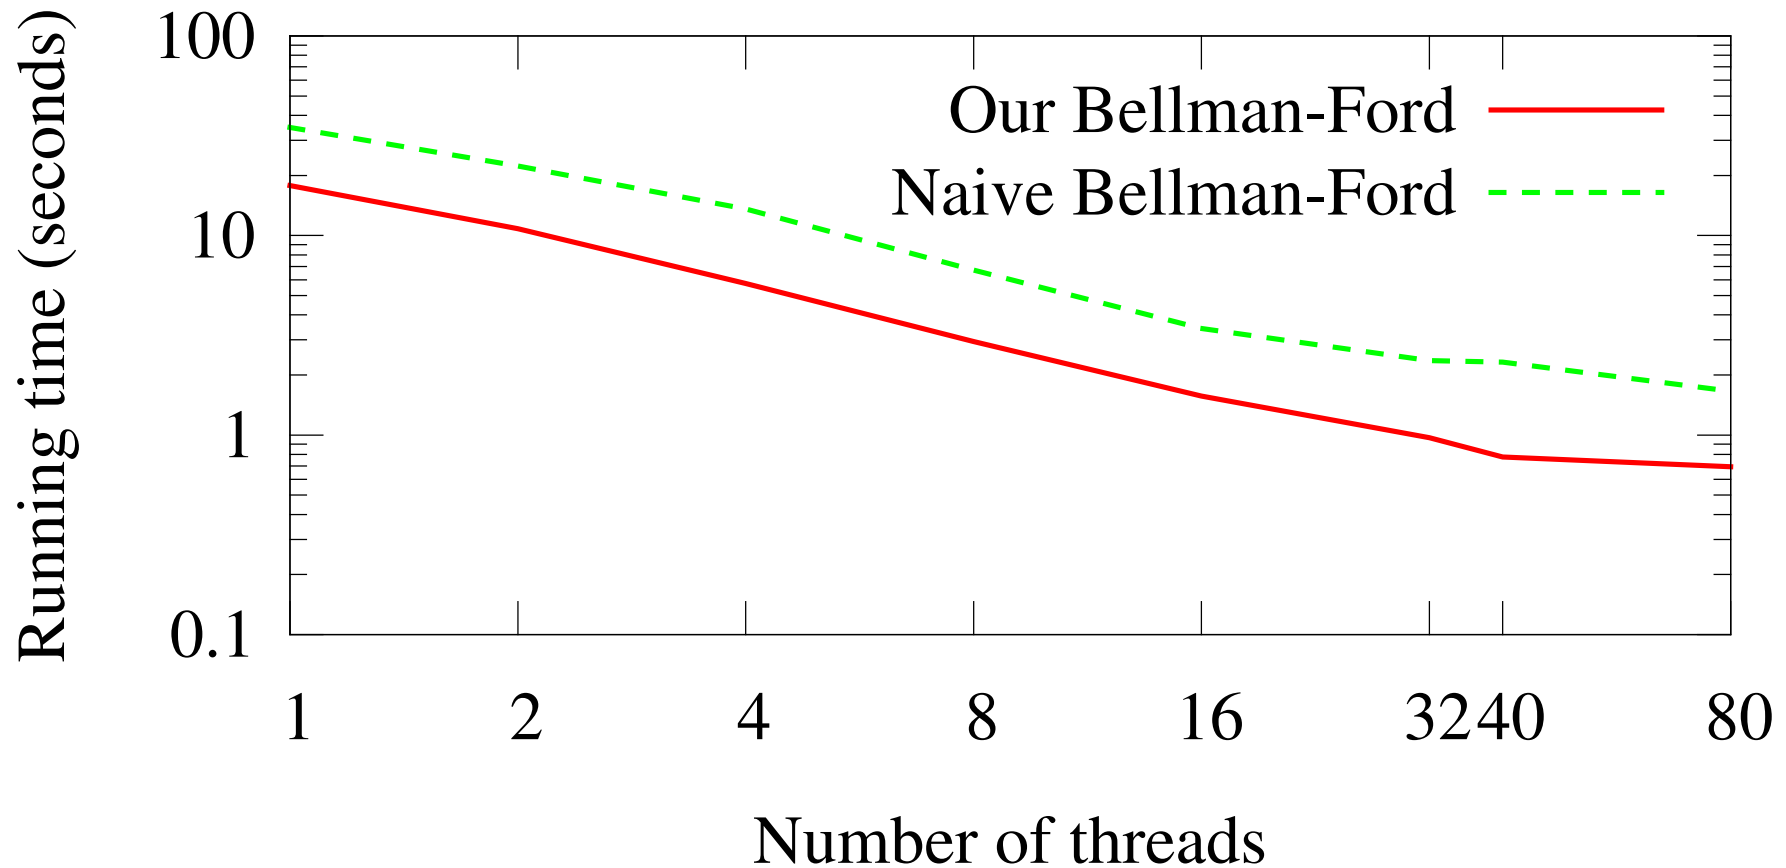


# Parallel Bellman-Ford Shortest Paths



# Parallel Bellman-Ford Performance

Times for Bellman-Ford on rMat24




# K-core decomposition

- A **k-core** of a graph is a maximal connected subgraph in which all vertices have degree at least  $k$
- A vertex has **core number**  $k$  if it belongs to a  $k$ -core but not a  $(k+1)$ -core
- Algorithm: Takes an unweighted, undirected graph  $G$  and returns the core number of each vertex in  $G$

```
k = 1
while(G is not empty) {
    while(there exists vertices with degree < k in G) {
        assign a core number of  $k-1$  to all vertices with degree <  $k$ ;
        remove all vertices with degree <  $k$  from  $G$ ;
    }
    k = k+1;
}
```

# PageRank



$$\text{PR}[v] = \frac{1 - \gamma}{|V|} + \gamma \sum_{u \in N^{-}(v)} \frac{\text{PR}[u]}{\text{deg}^{+}(u)}$$

# PageRank in Ligra

```
p_curr = {1/|V|, ..., 1/|V|};           p_next = {0, ..., 0};           diff = {};
```

```
procedure UPDATE(s, d):  
    return atomic_increment(p_next[d], p_curr[s] / degree(s));
```

```
procedure COMPUTE(i):  
    p_next[i] =  $\alpha \cdot p\_next[i] + (1 - \alpha) \cdot (1/|V|)$ ;  
    diff[i] = abs(p_next[i] - p_curr[i]);  
    p_curr[i] = 0;  
    return 1;
```

```
procedure PageRank(G,  $\alpha$ ,  $\epsilon$ ):
```

```
    frontier = {0, ..., |V|-1};
```

```
    error =  $\infty$ 
```

```
    while (error >  $\epsilon$ ):
```

```
        frontier = EDGEMAP(G, frontier, UPDATE, CONDtrue);
```

```
        VERTEXMAP(frontier, COMPUTE);
```

```
        error = sum of diff entries;
```

```
        swap(p_curr, p_next)
```

```
    return p_curr;
```

# PageRank

- *Sparse version?*
  - PageRank-Delta: Only update vertices whose PageRank value has changed by more than some  $\Delta$ -fraction (discussed in GraphLab and McSherry WWW '05)

# PageRank-Delta in Ligra

```
PR[i] = {1/|V|, ..., 1/|V|};
nghSum = {0, ..., 0};
Change = {}; //store changes in PageRank values


procedure UPDATE(s, d): //passed to EdgeMap
    return atomic_increment(nghSum[d], Change[s] / degree(s));

procedure COMPUTE(i): //now passed to VertexFilter
    Change[i] =  $\alpha \cdot \text{nghSum}[i]$ ;
    PR[i] = PR[i] + Change[i];
    return (abs(Change[i]) >  $\Delta$ ); //check if absolute value of change is big enough

procedure PageRank(G,  $\alpha$ ,  $\epsilon$ ):
    ...
    frontier = VERTEXFILTER(frontier, COMPUTE);
    ...
```

# Eccentricity estimation


- Takes an unweighted, undirected graph  $G = (V, E)$
- Returns an estimate of the eccentricity of each vertex where
- The **eccentricity** of a vertex  $v$  is the distance to furthest reachable vertex from  $v$





# Multiple BFS's

- Run multiple BFS's from a sample of random vertices and use distance from furthest sample as eccentricity estimate




$$\hat{ecc}(v) = \max(d(v, s_1), d(v, s_2), d(v, s_3), d(v, s_4)) \quad \text{for all } v$$

- In practice, need to run two sweeps to get good accuracy [KDD 2015]

# Eccentricity estimation implementation

- Run all  $k$  BFS's simultaneously
- Take advantage of bit-level parallelism



# Eccentricity estimation implementation

- Initial **frontier** =  $\{s_1, s_2, \dots, s_k\}$
  - $d = 0$
  - While **frontier** not empty:
    - **nextFrontier** =  $\{\}$
    - $d = d+1$
    - For each vertex  $v$  in **frontier**:
      - For each neighbor  $ngh$ :
        - Do bitwise-OR of  $v$ 's words with  $ngh$ 's words and store in  $ngh$
        - If any of  $ngh$ 's words changed:
          - $e\hat{c}(ngh) = \max(e\hat{c}(ngh), d)$  and place  $ngh$  on **nextFrontier** if not there
    - **frontier** = **nextFrontier**
- //Advance all BFS's by 1 level*
- atomic bitwise-OR using compare-and-swap
- EdgeMap
- //pass "visited" information*
- 

- We will implement this example in Ligra for  $k=64$

# Ligra Implementation Details

# VertexSubset, VertexMap, and VertexFilter


VertexSubset has one of two representations:

- Sparse integer array, e.g. {1, 4, 7}
- Dense boolean array, e.g. {0,1,0,0,1,0,0,1}

```
procedure VERTEXMAP(VertexSubset U, func F):  
  parallel foreach v in U:  
    F(v); //side-effects application data
```

```
procedure VERTEXFILTER(VertexSubset U, bool func F):  
  result = {}  
  parallel foreach v in U:  
    if(F(v) == 1) then:  
      add v to result;  
  return result;
```

# Sparse or Dense EdgeMap?

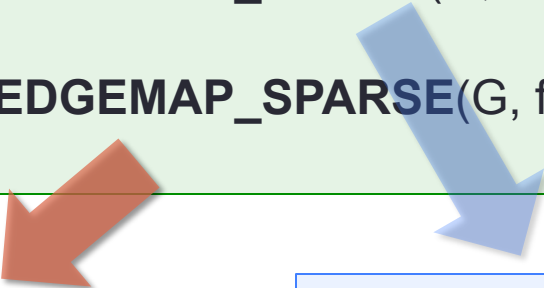


- Dense method better when frontier is large and many vertices have been visited
- Sparse (traditional) method better for small frontiers
- Switch between the two methods based on frontier size [Beamer et al. SC '12]

*Limited to BFS?*

# EdgeMap

```
procedure EDGEMAP(G, frontier, Update, Cond):  
  if (above threshold) then:  
    return EDGEMAP_DENSE(G, frontier, Update, Cond);  
  else:  
    return EDGEMAP_SPARSE(G, frontier, Update, Cond);
```



Loop through outgoing edges  
of frontier vertices in parallel

Loop through incoming edges of  
“unexplored” vertices (in parallel),  
breaking early if possible

- **More general than just BFS!**
- Generalized to many other problems
- Users need not worry about this

# EdgeMap (sparse version)

- How to represent VertexSubset?
  - Array of integers, e.g.  $U = \{0, 5, 7\}$

```

procedure EDGEMAP_SPARSE(G, frontier, UPDATE, COND):
  nextFrontier = {};
  parallel foreach v in frontier:
    parallel foreach w in out_neighbors(v):
      if(COND(w) == 1 and UPDATE(v, w) == 1) then:
        add w to nextFrontier;
  remove duplicates from nextFrontier;
  return nextFrontier;

```

```

parents = {-1, ..., -1};           // -1 indicates "unvisited"

procedure UPDATE(s, d):
  return compare_and_swap(parents[d], -1, s);

procedure COND(i):
  return parents[i] == -1; // checks if "unvisited"

procedure BFS(G, r):
  parents[r] = r;
  frontier = {r}; // vertexSubset
  while (size(frontier) > 0):
    frontier = EDGEMAP(G, frontier, UPDATE, COND);

```



# EdgeMap (dense version)

- How to represent dense VertexSubset?
  - Byte array, e.g.  $U = \{1, 0, 0, 0, 0, 1, 0, 1\}$

procedure **EDGEMAP\_DENSE**(G, frontier, **UPDATE**, **COND**):

  nextFrontier = {0, ..., 0};

  parallel foreach v in G:

    if (**COND**(v) == 1) then:

      foreach ngh in in\_neighbors(v):

        if ngh in frontier and **UPDATE**(ngh, v) == 1 then:

          nextFrontier[v] = 1;

        if (**COND**(v) == 0) then:

          break;

  return nextFrontier;

parents = {-1, ..., -1}; *// -1 indicates "unvisited"*

procedure **UPDATE**(s, d):

  return compare\_and\_swap(parents[d], -1, s);

procedure **COND**(i):

  return parents[i] == -1; *// checks if "unvisited"*

procedure **BFS**(G, r):

  parents[r] = r;


  frontier = {r}; *// vertexSubset*

  while (size(frontier) > 0):

    frontier = **EDGEMAP**(G, frontier, **UPDATE**, **COND**);



# Frontier-based approach enables sparse/dense traversal

Twitter graph (41M vertices, 1.5B edges)



# Ligra BFS Performance


Twitter graph (41M vertices, 1.5B edges)



- Comparing against direction-optimizing code by Beamer et al.

# Ligra PageRank Performance

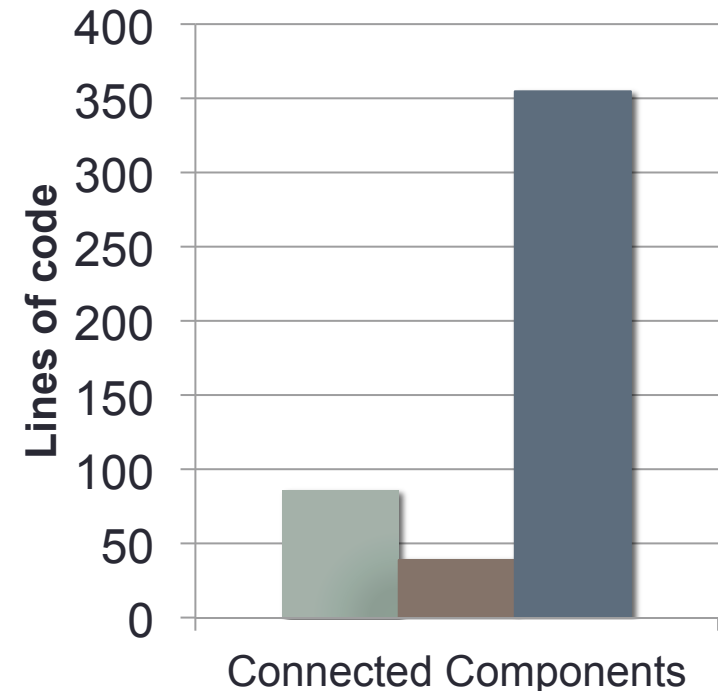

Twitter graph (41M vertices, 1.5B edges)



- Easy to implement “sparse” version of PageRank in Ligra



# Ligra Connected Components Performance

Twitter graph (41M vertices, 1.5B edges)



- Performance close to hand-written code
- Faster than existing high-level frameworks at the time
- Shared-memory graph processing is very efficient
  - Several shared-memory graph processing systems subsequently developed: Galois [SOSP '13], X-stream [SOSP '13], PRISM [SPAA '14], Polymer [PPoPP '15], Ringo [SIGMOD '15], GraphMat [VLDB '15]


# Large Graphs



- Most can fit on commodity shared memory machine
- *What if you don't have that much memory?*

# Ligra+: Adding Graph Compression to Ligra


# Ligra+: Adding Graph Compression to Ligra



- Same interface as Ligra
- All changes hidden from the user!




# Graph representation



- Graph reordering to improve locality
  - Goal: give neighbors IDs close to vertex ID
  - BFS, DFS, METIS, our own separator-based algorithm

# Variable-length codes

- k-bit codes
  - Encode value in chunks of k bits
  - Use k-1 bits for data, and 1 bit as the “continue” bit
- Example: encode “401” using 8-bit (byte) code



# Encoding optimization

- Another idea: get rid of “continue” bits



Number of bytes  
required to encode  
each integer

1

2

2

2

2

2

2

2

.....

Use run-length encoding

Header




Integers in group  
encoded in byte chunks

Number of bytes  
per integer

Size of group  
(max 64)

- Increases space, but makes decoding cheaper (no branch misprediction from checking “continue” bit)

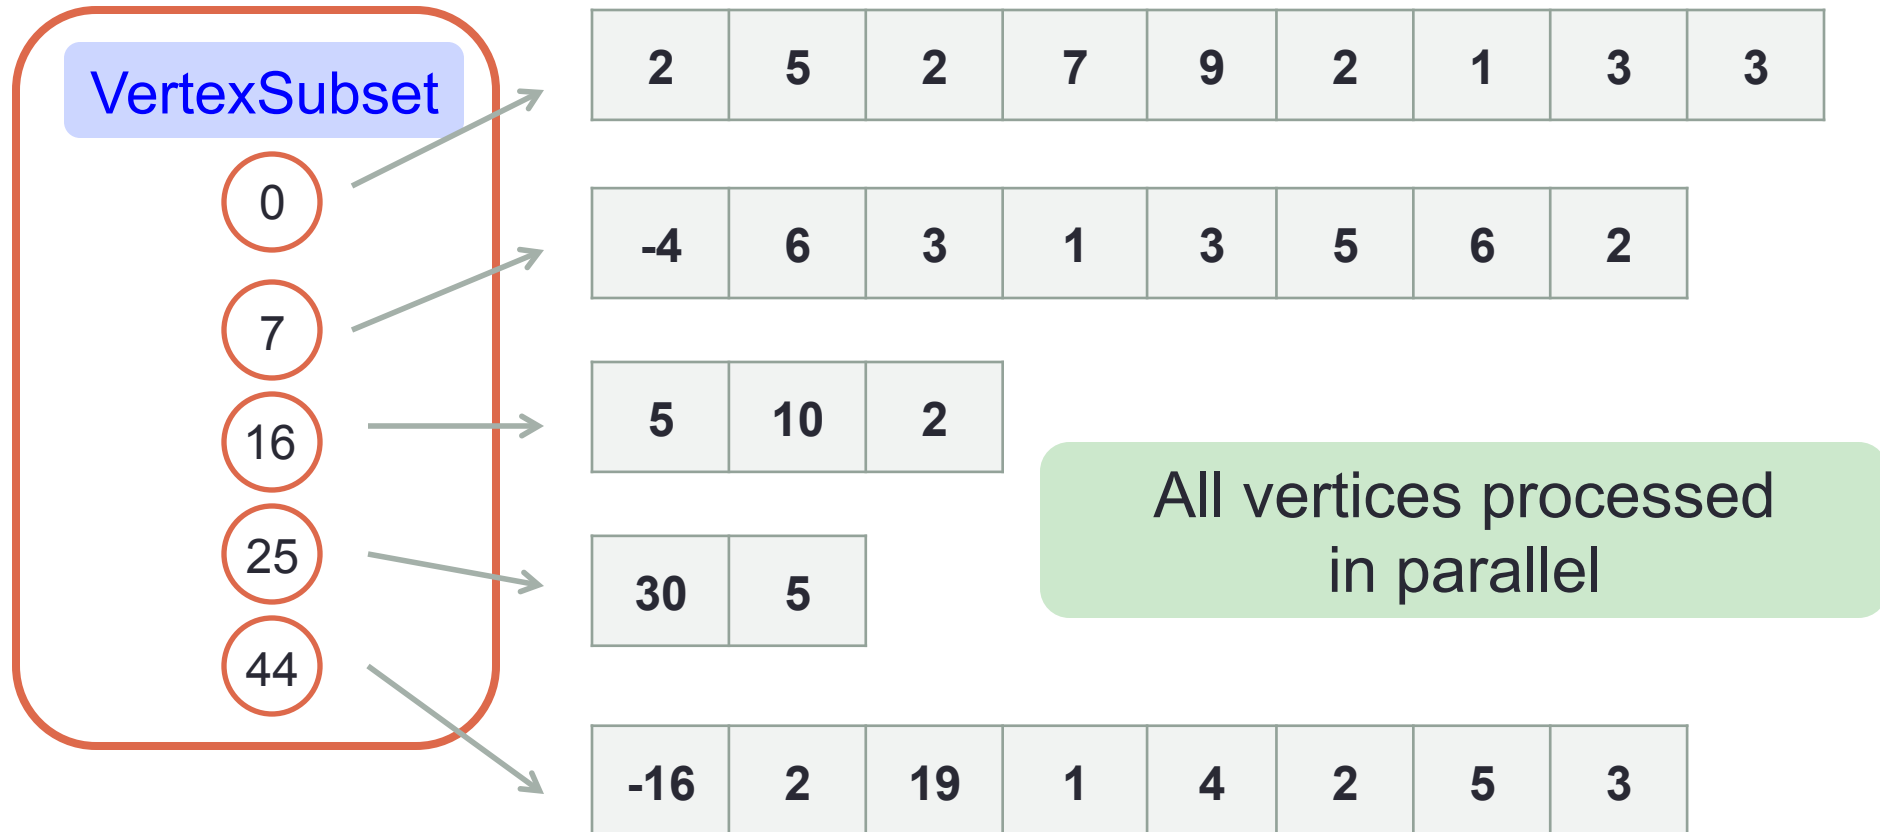
# Ligra+: Adding Graph Compression to Ligra



- Same interface as Ligra
- All changes hidden from the user!

# Modifying EdgeMap

- Processes outgoing edges of a subset of vertices



*What about high-degree vertices?*

# Handling high-degree vertices

High-degree  
vertex

-1	2	4	3	16	2	1	5	8	19	4	1	23	14	12	1	9	10	3	5	...
----	---	---	---	----	---	---	---	---	----	---	---	----	----	----	---	---	----	---	---	-----

Chunks of size  $T$

-1	2	4	3	16	2	27	5	8	19	4	1	87	14	12	1	9	10	...
----	---	---	---	----	---	----	---	---	----	---	---	----	----	----	---	---	----	-----


Encode first entry relative to source vertex

All chunks can be  
decoded in parallel!


- We chose  $T=1000$
- Similar performance and space usage for a wide range of  $T$

# Ligra+: Adding Graph Compression to Ligra

Space relative to Ligra



40-core time relative to Ligra



- Using 8-bit codes with run-length encoding
- Cost of decoding on-the-fly?
- Memory bottleneck a bigger issue as graph algorithms are memory-bound

# Demo on compressed graphs in Ligra



# Other Graph Processing Systems

# Many existing frameworks


- Pregel/Giraph/GPS, GraphLab/PowerGraph, PRISM, Pegasus, Knowledge Discovery Toolbox/CombBLAS, GraphChi, GraphX, Galois, X-Stream, Gunrock, GraphMat, Ringo, TurboGraph, FlashGraph, Grace, PathGraph, Polymer, GoFFish, Blogel, LightGraph, MapGraph, PowerLyra, PowerSwitch, XDGP, Signal/Collect, PrefEdge, Parallel BGL, KLA, Grappa, Chronos, Green-Marl, GraphHP, P++, LLAMA, Venus, Cyclops, Medusa, NScale, Neo4J, Trinity, GBase, HyperGraphDB, Horton, GSPARQL, Titan, and many others...
- Cannot list everything here. For more information, see:
  - “Systems for Big Graphs”, [Khan and Elnikety VLDB 2014 Tutorial](#)
  - “Trade-offs in Large Graph Processing: Representations, Storage, Systems and Algorithms”, [Ajwani et al. WWW 2015 Tutorial](#)
  - “A Survey of Parallel Graph Processing Frameworks”, [Doekemeijer and Varbanescu 2014](#)
  - “Thinking like a Vertex: A Survey of Vertex-Centric Frameworks for Large-Scale Distributed Graph Processing”, [McCune et al. 2015](#)

# Pregel

- “Think like a vertex”
- Distributed-memory, uses message passing
- Bulk synchronous model

```
template <typename VertexValue,  
          typename EdgeValue,  
          typename MessageValue>  
class Vertex {  
public:  
    virtual void Compute(MessageIterator* msgs) = 0;  
  
    const string& vertex_id() const;  
    int64 superstep() const;  
  
    const VertexValue& GetValue();  
    VertexValue* MutableValue();  
    OutEdgeIterator GetOutEdgeIterator();  
  
    void SendMessageTo(const string& dest_vertex,  
                      const MessageValue& message);  
    void VoteToHalt();  
};
```

- Vertices can be either “active” or “inactive”




# GraphLab/PowerGraph

- “Think like a vertex”
- Vertices define Gather, Apply, and Scatter functions
- Data on edges as well as vertices
- Scheduler processes “active” vertices
  - Supports asynchronous execution
    - Useful for some machine learning applications
  - Different levels of consistency
  - Different scheduling orders
- Current version for distributed memory (original version was for shared memory)

```
interface GASVertexProgram(u) {
  // Run on gather_nbrs(u)
  gather( $D_u$ ,  $D_{(u,v)}$ ,  $D_v$ )  $\rightarrow$  Accum
  sum(Accum left, Accum right)  $\rightarrow$  Accum
  apply( $D_u$ , Accum)  $\rightarrow D_u^{\text{new}}$ 
  // Run on scatter_nbrs(u)
  scatter( $D_u^{\text{new}}$ ,  $D_{(u,v)}$ ,  $D_v$ )  $\rightarrow (D_{(u,v)}^{\text{new}}, \text{Accum})$ 
}
```

Vertex Data: 

Edge Data: 




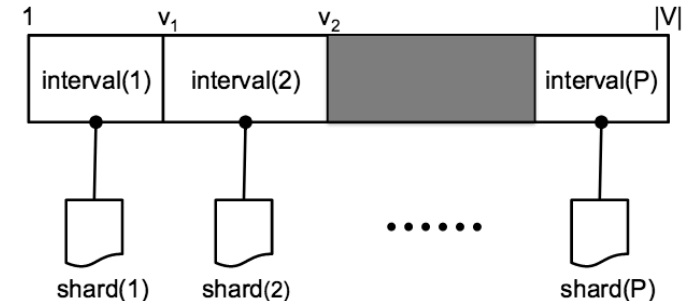
GraphLab: A New Framework for Parallel Machine Learning, Low et al. UAI 2010

PowerGraph: Distributed Graph-Parallel Computation on Natural Graphs, Gonzalez et al.


OSDI 2012

# GraphChi

- “Think like a vertex”
  - Define Update function for vertices
- Optimized for disk-based execution
- Divides edges into “shards”
  - Each shard has a range of target IDs
  - Each shard sorted by source ID
- Parallel sliding windows method for efficient execution
  - One shard in memory, other shards read in streaming fashion
  - About  $O((V+E)/B)$  I/O's per iteration



# Linear algebra abstraction



- PEGASUS: Framework using matrix-vector abstraction in MapReduce [Kang et al. ICDM 2009]
- Knowledge Discovery Toolbox/CombBLAS: Uses matrix-vector and matrix-matrix routines for implementing graph algorithms; uses Python frontend, C++/MPI backend [Lugowski et al. SDM 2012]

# Some other shared memory systems

- X-Stream [SOSP 2013]
  - Edge-centric abstraction
  - Supports disk-based execution
- Galois [SOSP 2013]
  - Parallel programming framework based on dynamic sets with various schedulers
  - Implements graph abstractions of Ligra, X-Stream, and GraphLab
- PRISM [SPAA 2014]
  - Deterministic version of GraphLab using graph coloring for scheduling; implemented in Cilk
- Polymer [PPoPP 2015]
  - NUMA-aware implementation of Ligra's abstraction
- GraphMat [VLDB 2015]
  - Uses optimized matrix-vector routines to implement graph algorithms
- Gunrock [PPoPP 2016]
  - Abstraction for frontier-based computations on GPUs

# Exercises



# Exercises

- Implement a version of BFS that gives a **deterministic** BFS tree (i.e., the Parents array is deterministic)
- Implement a version of BFS that uses a **bit-vector** to check the “visited” status before updating the Parents array



# Thank you!

Code: <https://github.com/jshun/ligra/>

## References

*Ligra: A Lightweight Graph Processing Framework for Shared Memory*, PPOPP 2013.

*Smaller and Faster: Parallel Processing of Compressed Graphs with Ligra+*, DCC 2015.

*An Evaluation of Parallel Eccentricity Estimation Algorithms on Undirected Real-World Graphs*, KDD 2015.